

# COMP2300-COMP6300-ENGN2219

## Computer Organization & Program Execution

Convener: Shoaib Akram  
[shoaib.akram@anu.edu.au](mailto:shoaib.akram@anu.edu.au)



Australian  
National  
University

# Recall: Temporal Parallelism



- Temporal Parallelism (pipelining):
  - Break down a circuit into stages
  - Each task passes through all stages
  - Multiple tasks are spread through stages

# Recall: Automotive Pipeline



# Recall: Pipelining

- If a task of latency  $L$  is broken into  $N$  stages, and all stages are of equal length, then the throughput is  $N/L$




- The challenge of pipelining is to find stages of equal length
- Let's go back to baking cookies

# Recall: Cookie Parallelism

- Ben and Jon are making cookies. Let's study the latency and throughput of rolling and baking many cookie trays with
  - No parallelism
  - Spatial parallelism
  - Pipelining
  - Spatial parallelism + pipelining


# No Parallelism (Ben Only)



Latency (**hours**/tray):

Throughput (**trays**/hour):

# Spatial Parallelism (Ben & Jon)

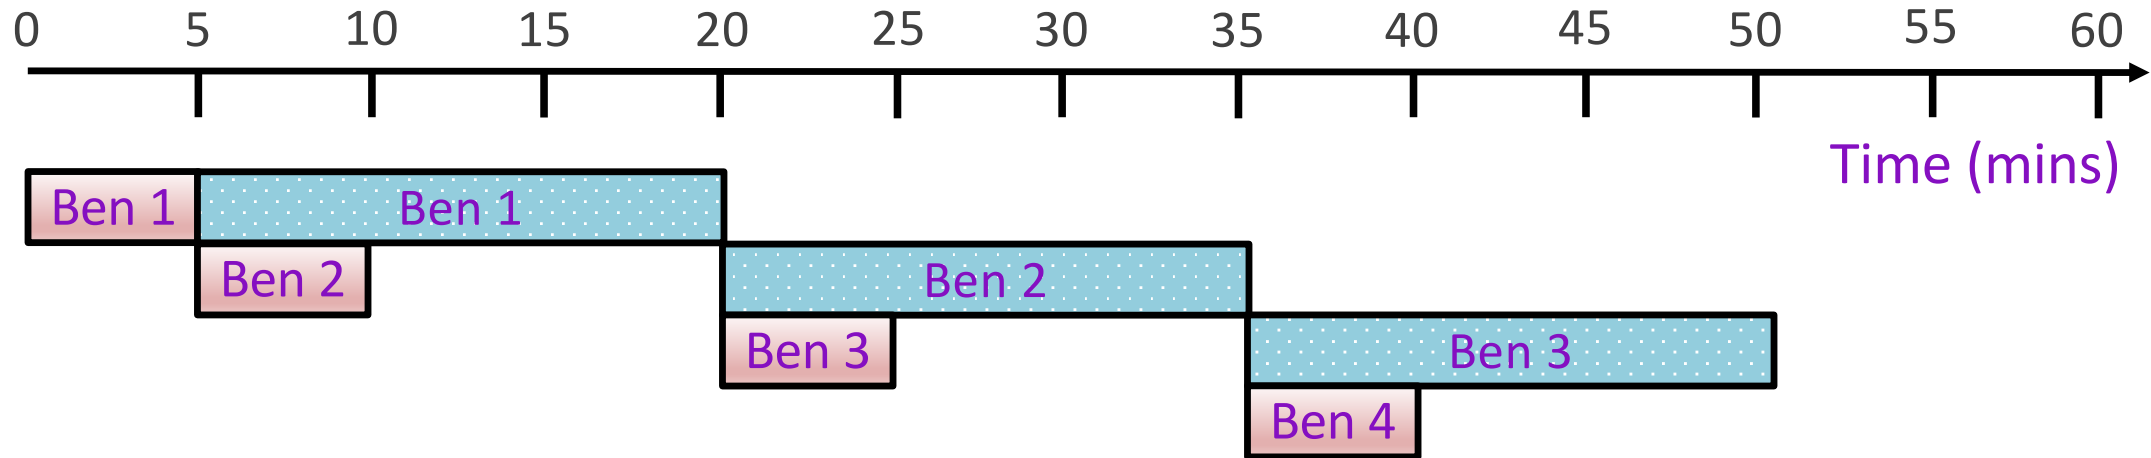


Latency (**hours**/tray):

Throughput (**trays**/hour):

Note: Jon owns a tray and oven (hardware duplication)

# Pipelining (Ben Only)




Latency (**hours**/tray):

Throughput (**trays**/hour):

Note: Ben decides not to waste a separate tray and oven



# Spatial + Temporal Parallelism



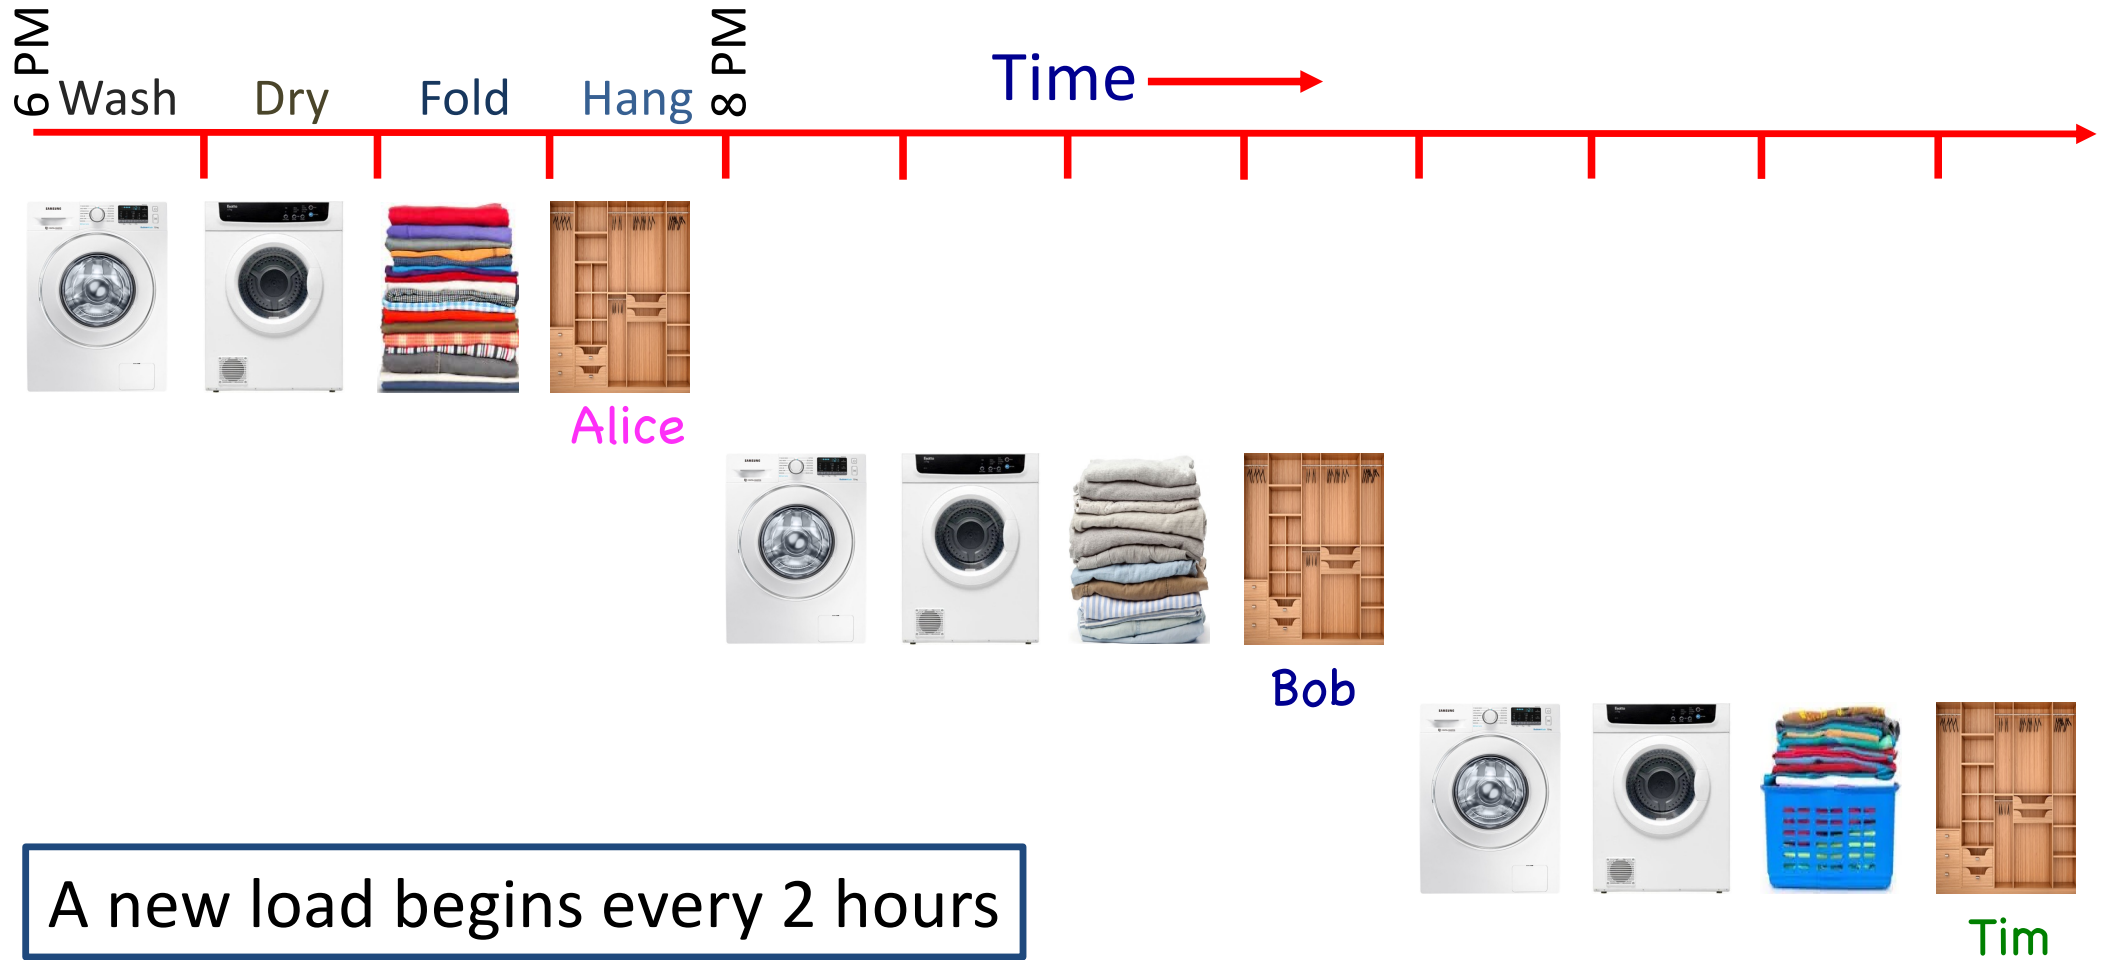
Latency (**hours**/tray):

Throughput (**trays**/hour):

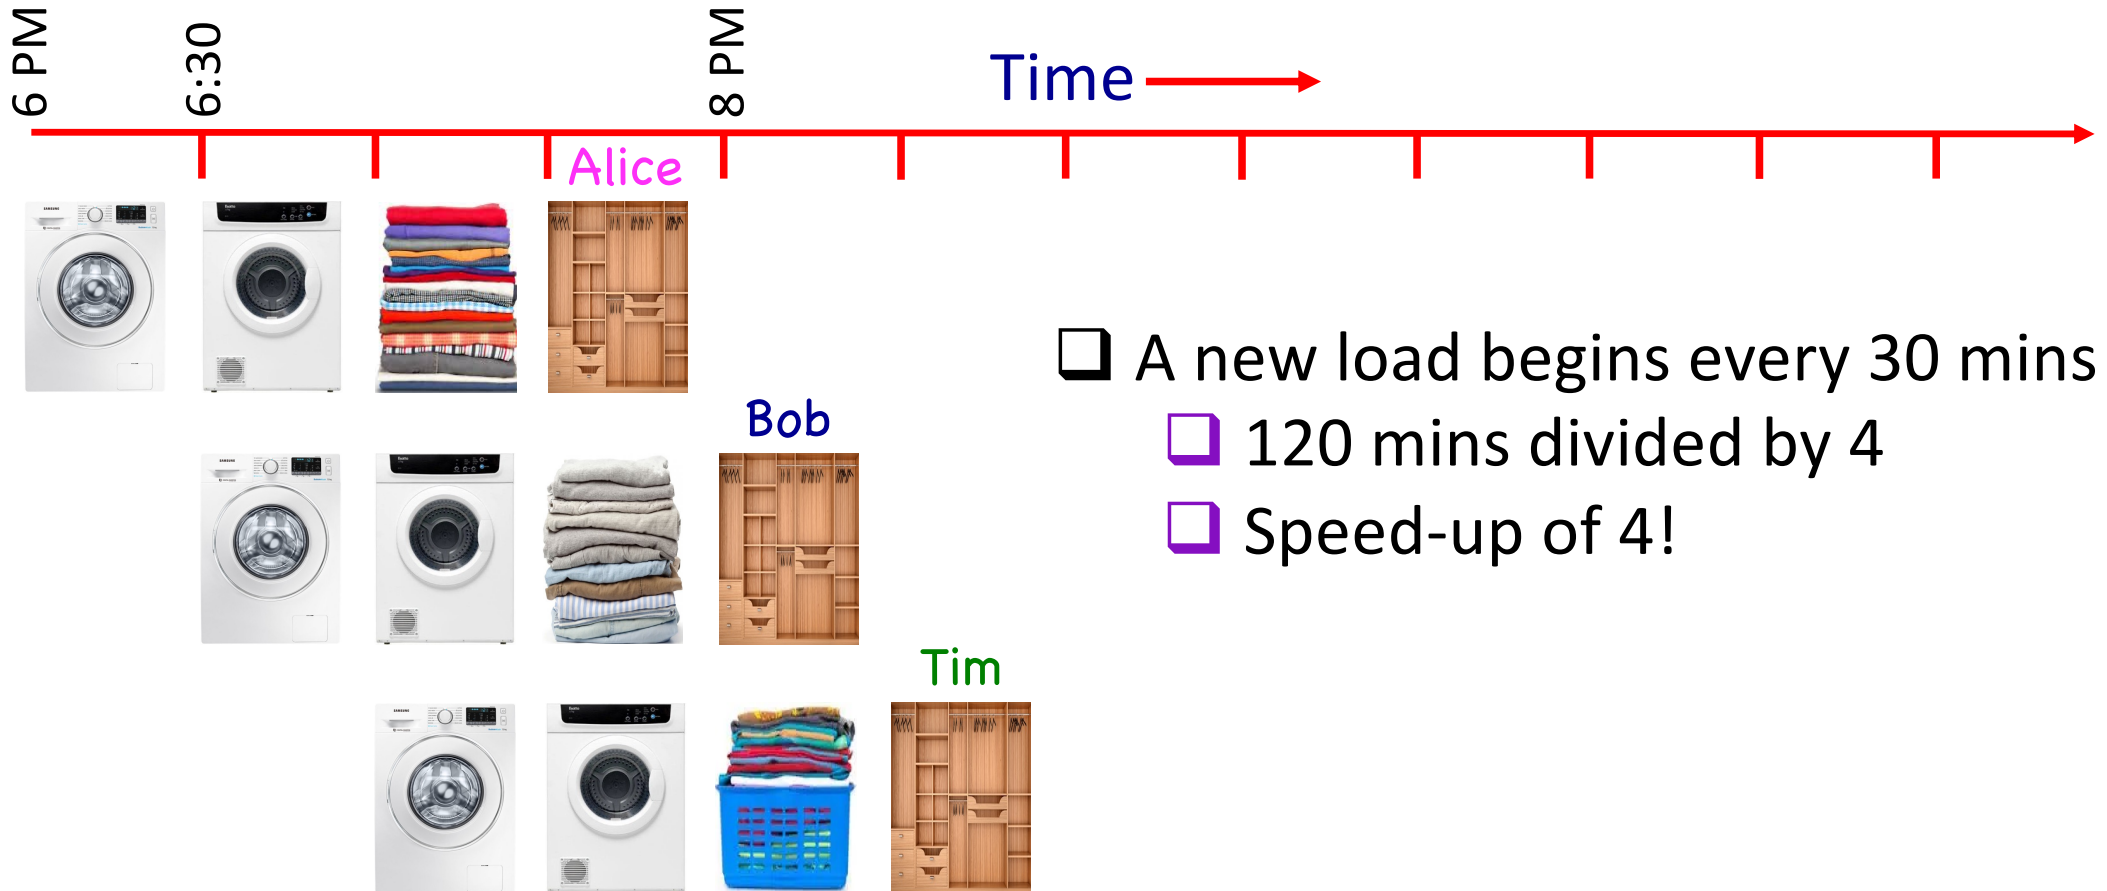
# Answers Explained

- **No parallelism**
  - Latency is clearly 20 minutes ( $1/3$  hours/tray)
  - Throughput is 3 trays per hour
- **Spatial parallelism**
  - Latency remains unchanged as it still takes 20 mins to finish a tray
  - Throughput is doubled via duplication: 6 trays per hour
- **Pipelining**
  - Latency for a single tray remains unchanged
  - Throughput: Ben puts a new tray in the oven every 15 minutes, so the throughput is 4 trays per hour
  - Note that in the first hour, Ben loses 5 minutes to fill the pipeline
- **Spatial parallelism + pipelining**
  - Latency remains unchanged
  - Throughput: Ben & Jon combo puts two trays in the oven every 15 minutes, so the throughput is 8 trays per hour

# Sequential Laundry



# Pipelined Laundry

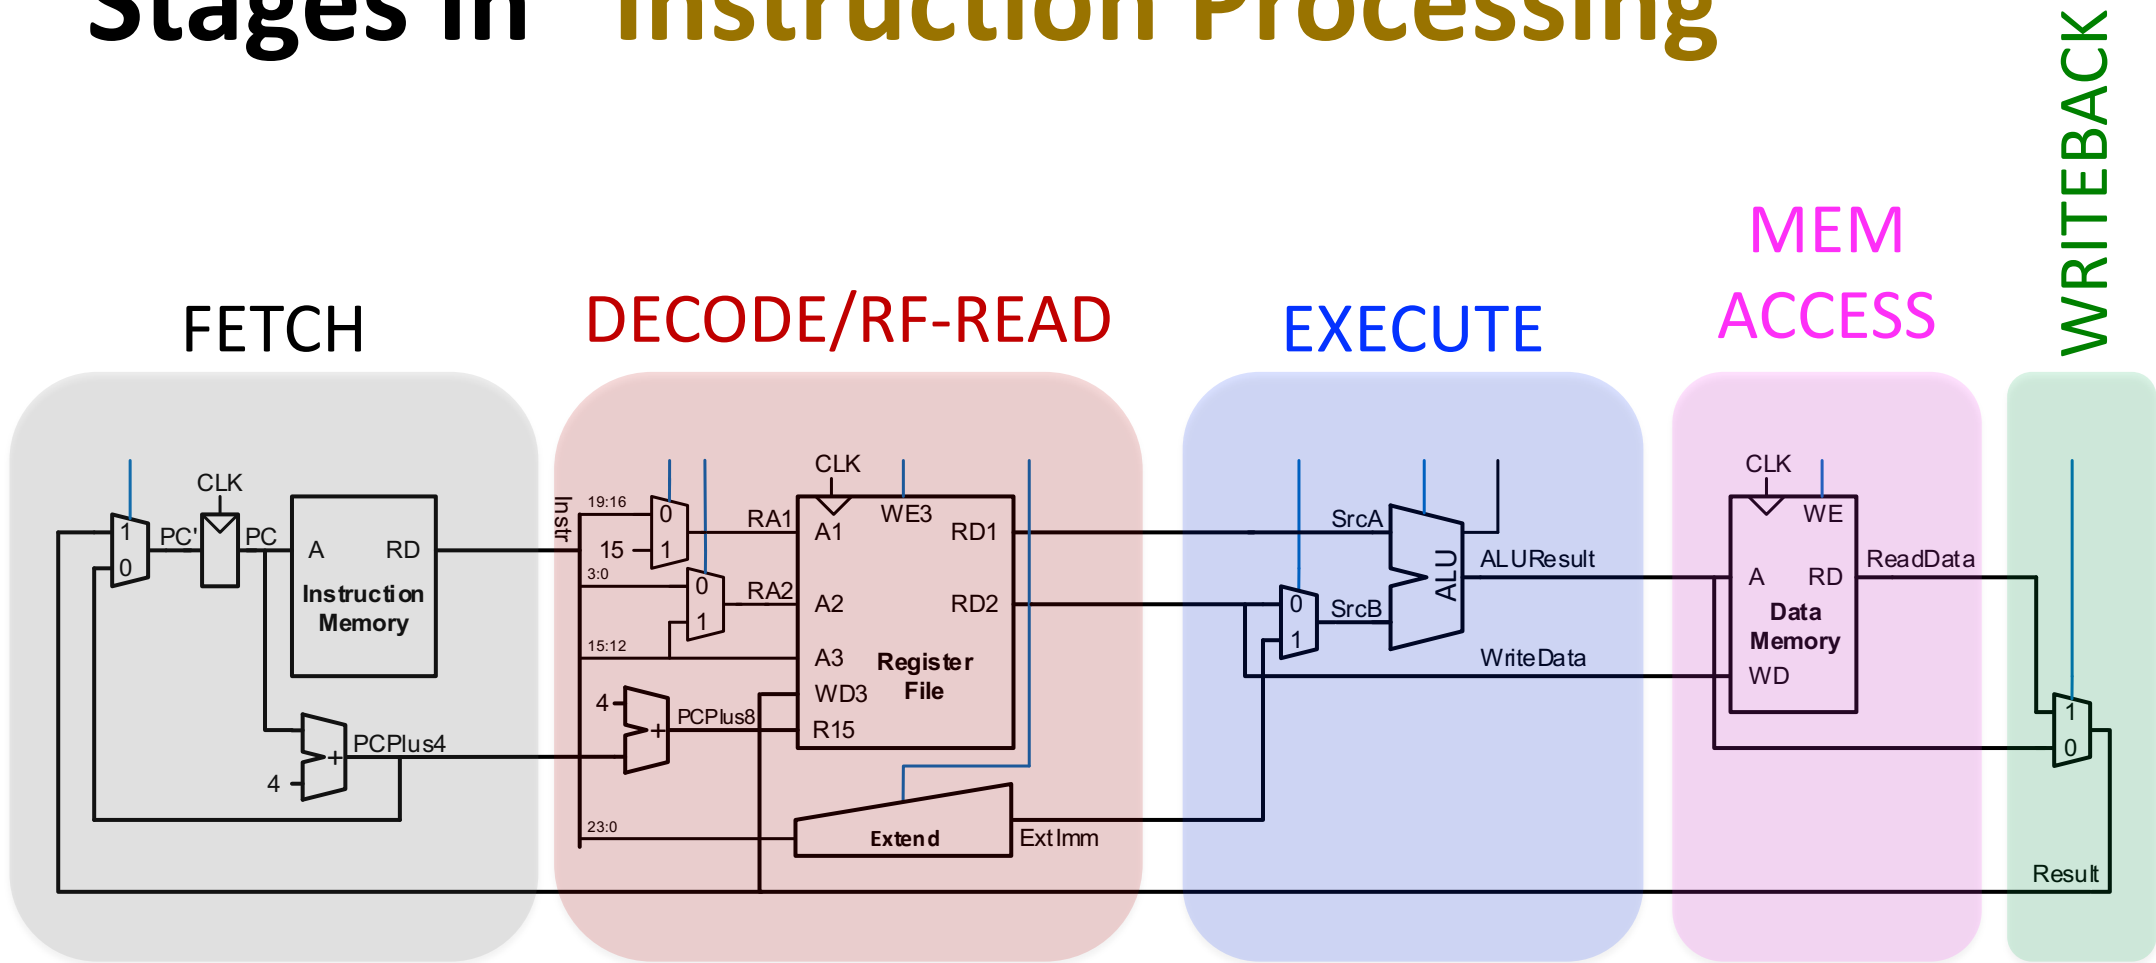


# Recall: **Pipelining** Circuits

- Divide a **large** combinational circuit into shorter **stages**
- Insert **registers** between the stages
  - The outputs of one stage are copied into a register and communicated to the next stage
- Run the **pipelined** circuit at a **higher** clock frequency
  - Each clock cycle, data flows through the pipeline from left to the right
  - Multiple tasks can be spread across the pipeline

# Pipelined Microarchitecture

# Stages in “Instruction Processing”



# Pipelined Microarchitecture: **Key Idea**

- **Multiple instructions** (up to 5) can be in the pipeline in any cycle
- **Each instruction can be in a different stage**
  - Idea is for “**maximizing utilization**” of hardware resources
- Stages must be isolated from one another using pipelined register (non-arch. registers). Referred to as “**PPR**”
- The work of a stage should be preserved in a **PPR** each cycle



# Key Idea (Continued)

- The work of a stage should be preserved in a PPR each cycle
- PPR acts as a source of data the next stage needs in a subsequent cycle
- If any subsequent stage down the pipeline needs data from an earlier stage it must be passed through the PPRs
  - .... Things don't always go smoothly as we shall see!

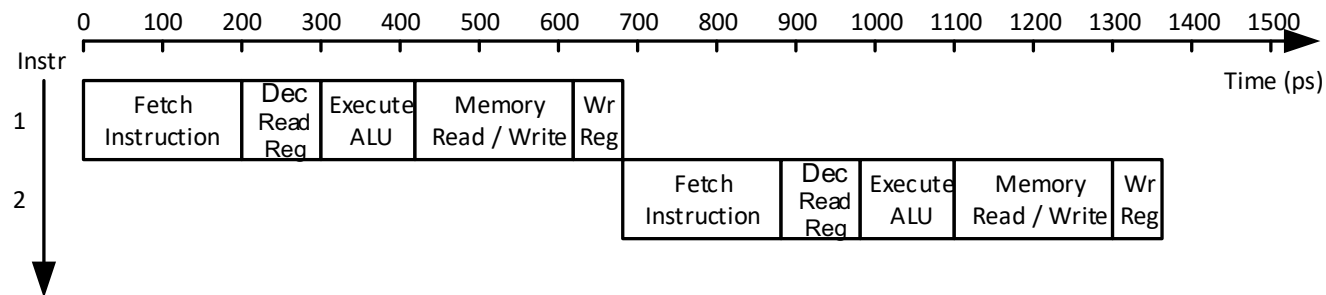
# Timing Diagrams

- To visualize the execution of many instructions in a pipeline we can use timing diagrams where:
  - Time is on the horizontal axis
  - Instructions are on the vertical axis

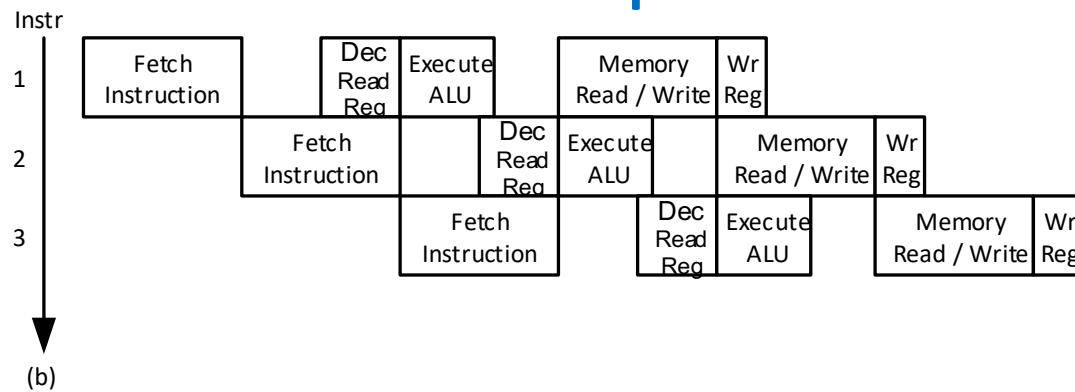
# Timing Diagrams

Assumption of logic element delays from Table 7.5 of textbook

## Single-Cycle



## Pipelined



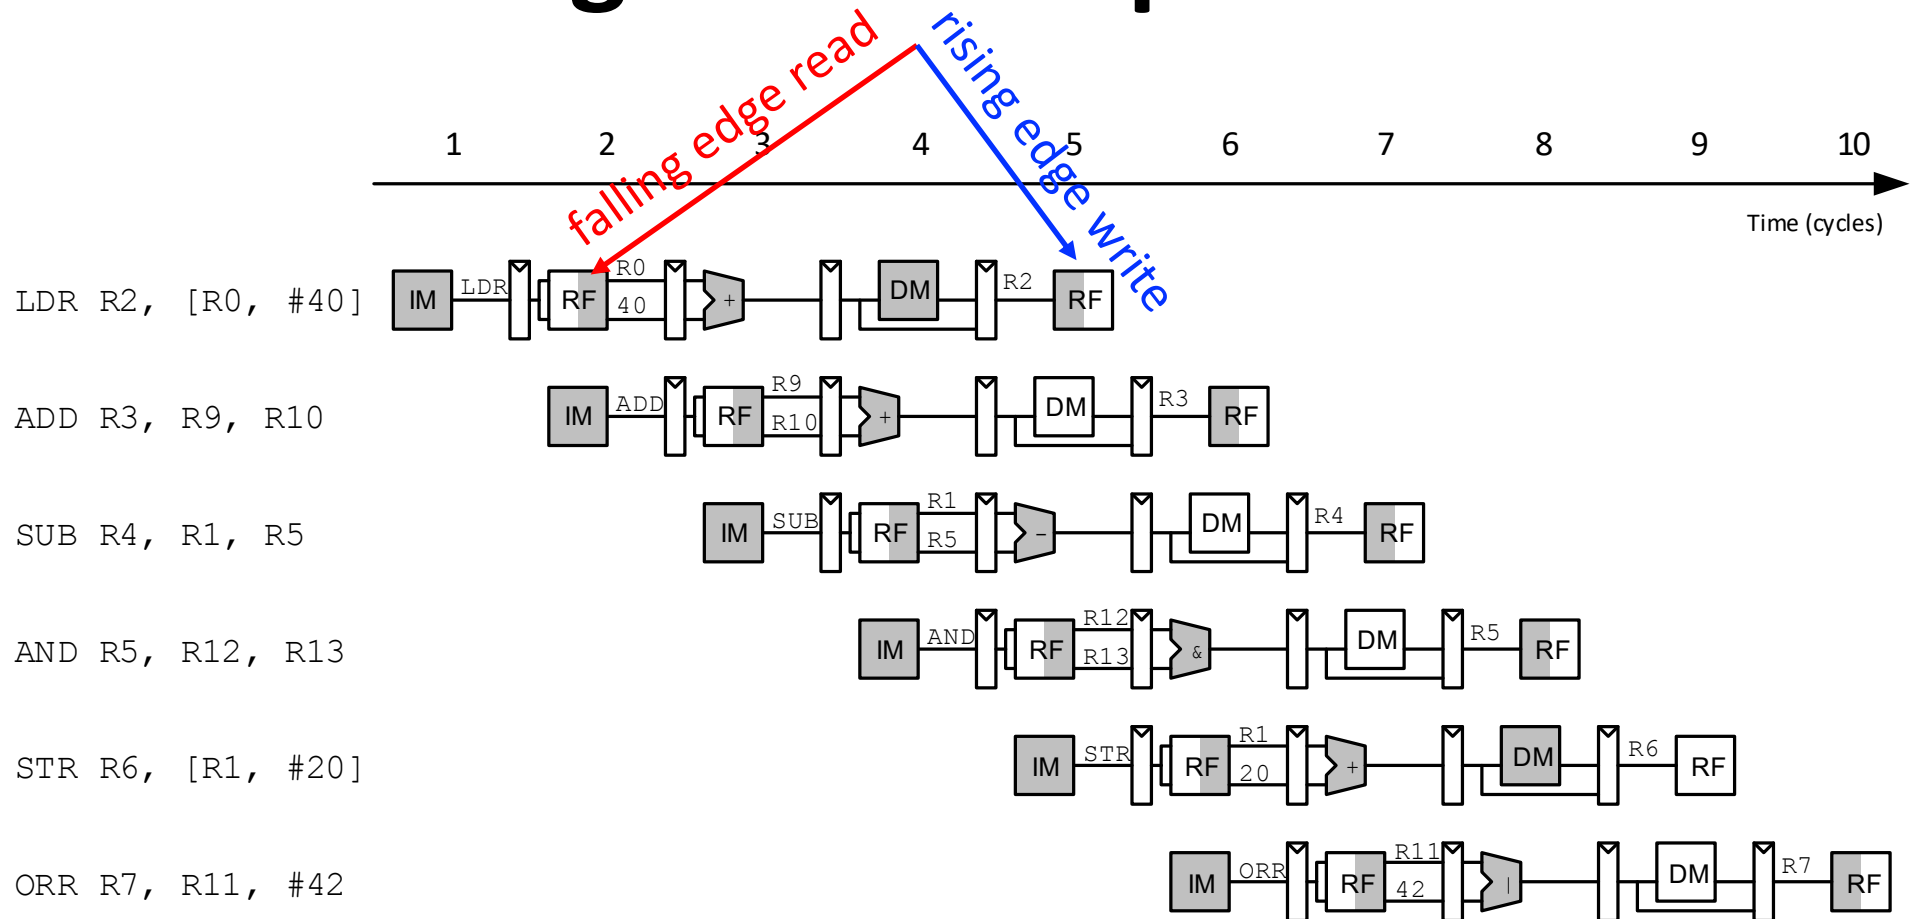
# Performance Analysis

- In the previous slide, what is the execution time and instructions per second (IPS) for the single-cycle microarchitecture?
  - 1.47 Billion Instructions per Second
- What about the pipelined microarchitecture?
  - The length of the pipeline stage is set by the slowest stage to be 200 ps
  - 1 instruction per 200 ps
  - 5 billion instructions per second

# Instruction Latency with Pipelining

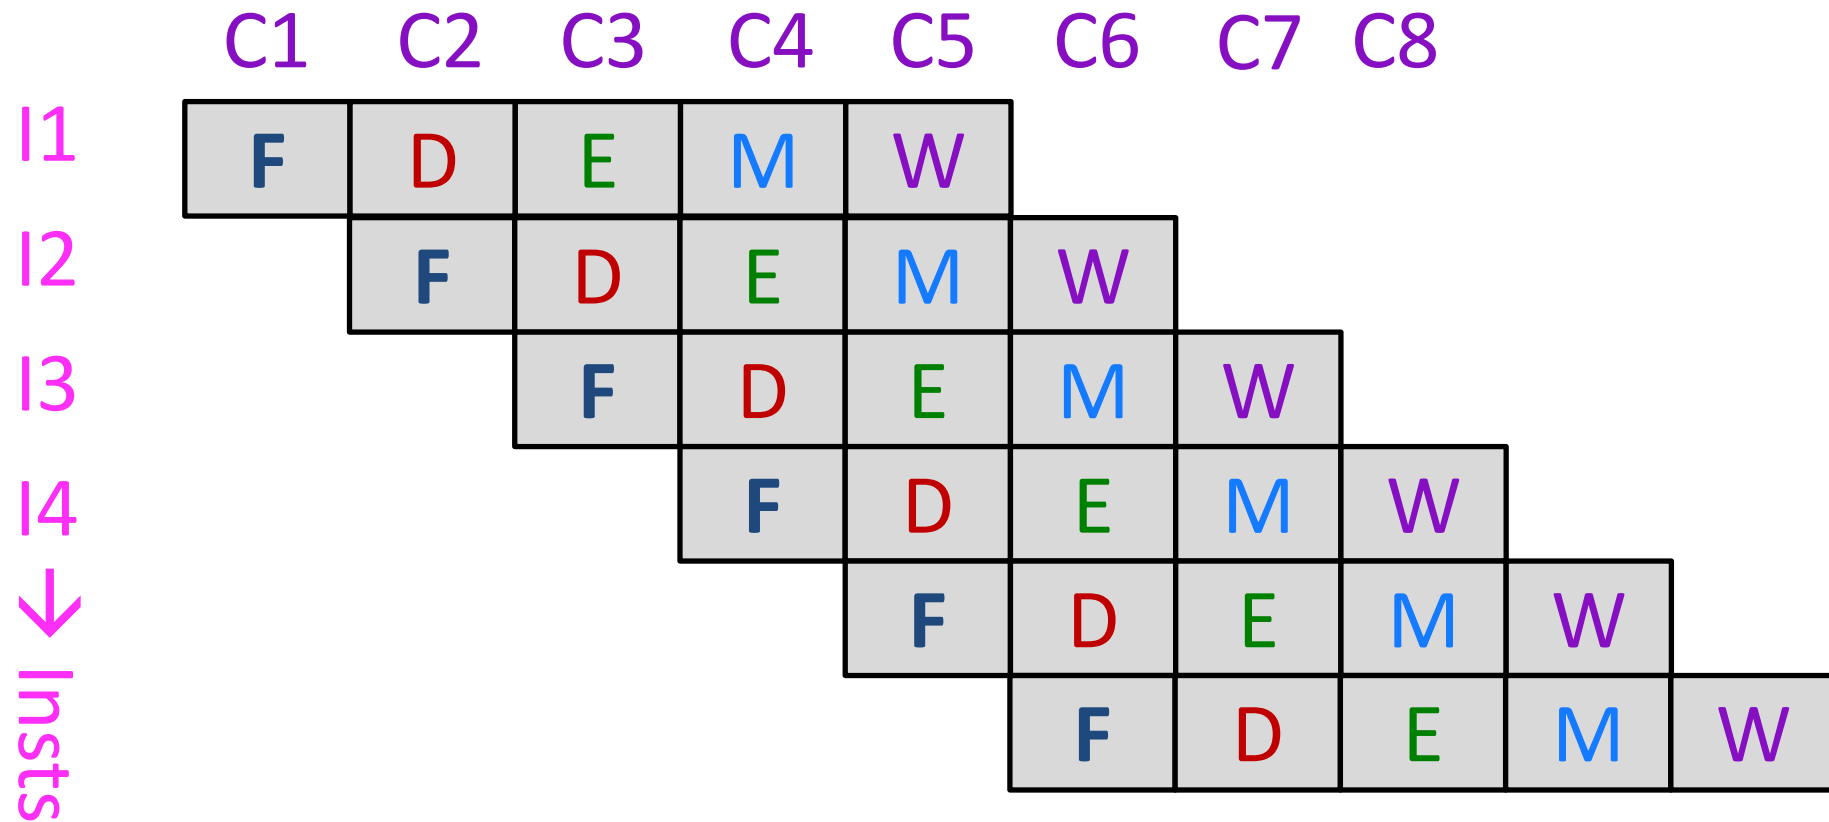
- Pipelining does not help to reduce the latency of a single instruction
- Latency of a single instruction increases
  - Sequencing overhead of pipeline registers
  - Clock cycle time decided by slowest pipeline stage (internal fragmentation due to imbalanced stages)
- Pipelining helps increase the throughput of an entire workload
  - Workload = Number of instructions
  - Workload must be “sufficiently” large

# Abstract Diagrams of Pipelined uArch



**Key idea:** In one cycle, an instruction writeback can be visible to a younger instruction's reg read

# Simplified View of Pipelining

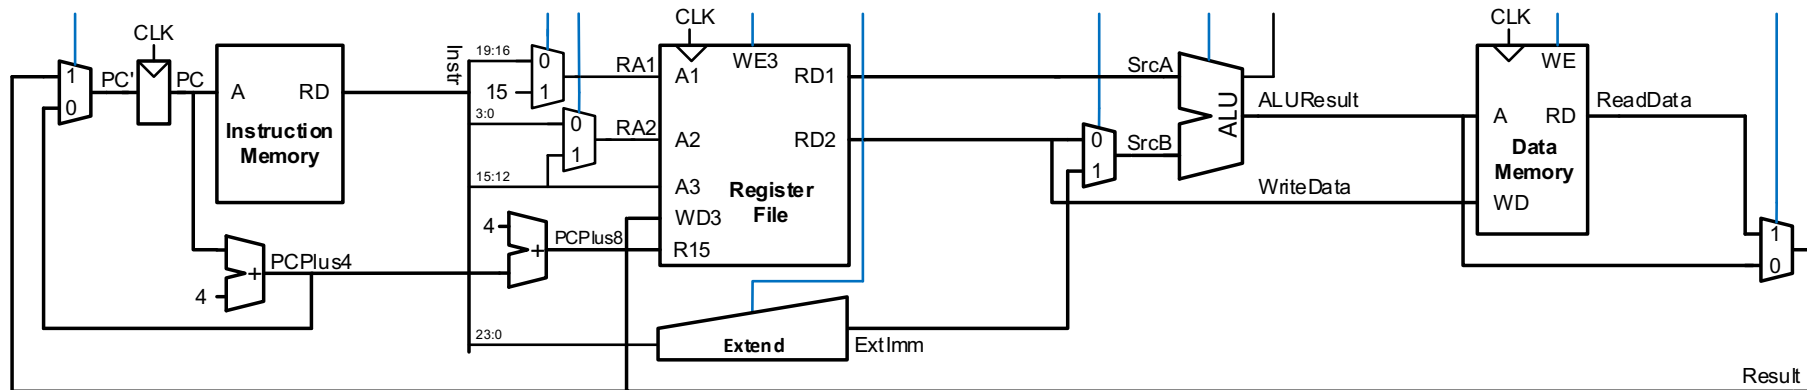


# Let's complete the picture

- Start with the single-cycle microarchitecture
- And insert pipeline registers



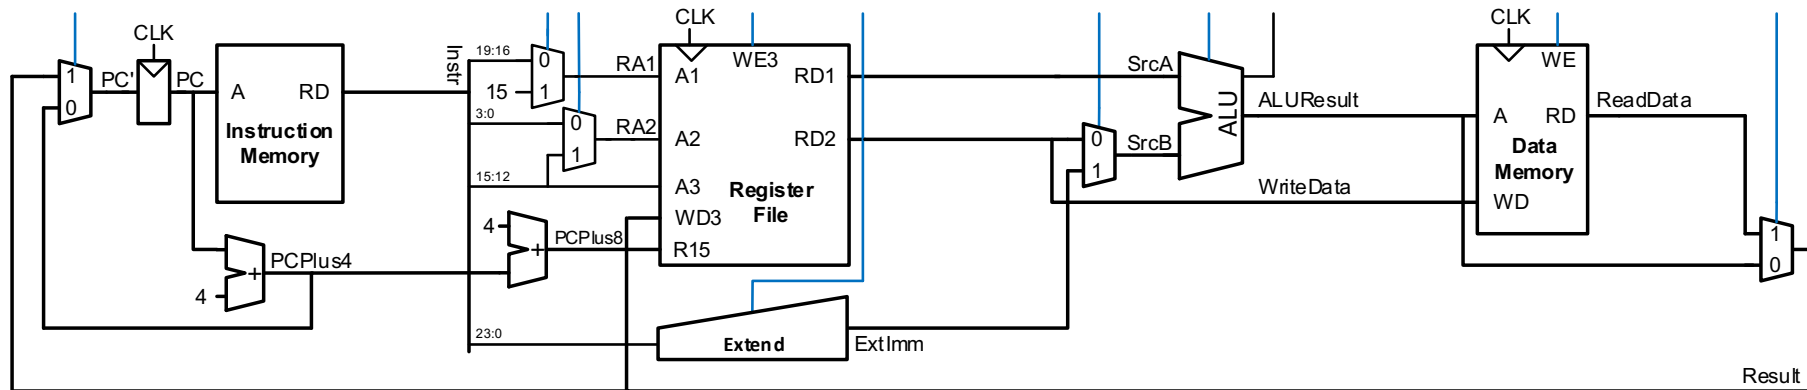
## Single-Cycle



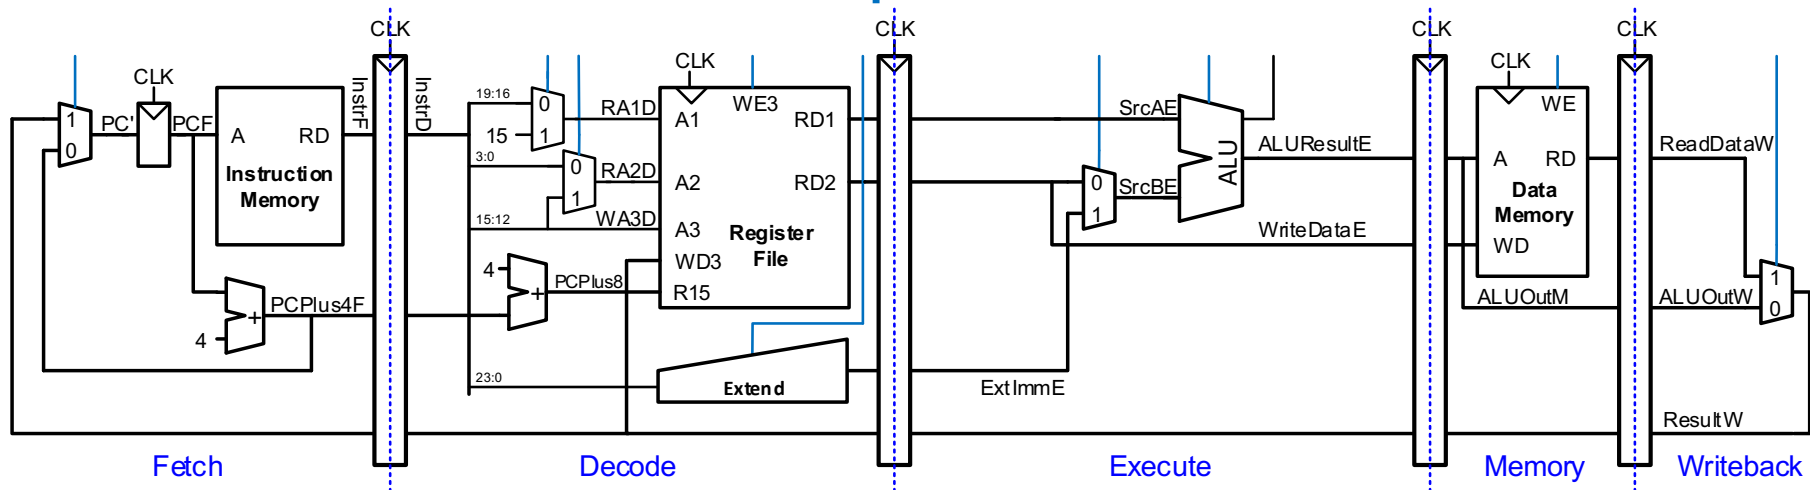
- Once we insert pipeline registers, we would need to pass the results of one stage to the next stage via the pipeline registers
- What is the outcome of the **FETCH** stage?

# Pipeline Microarchitecture

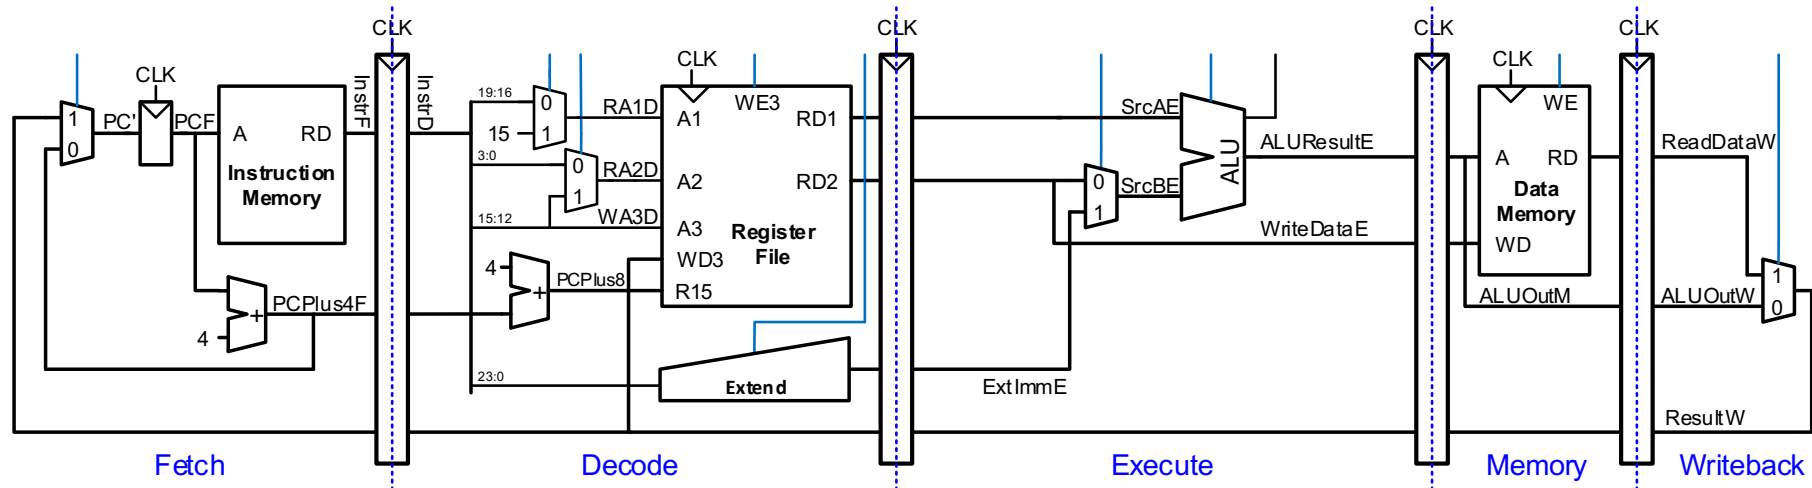
## Single-Cycle



## Pipelined



# Pipeline Microarchitecture



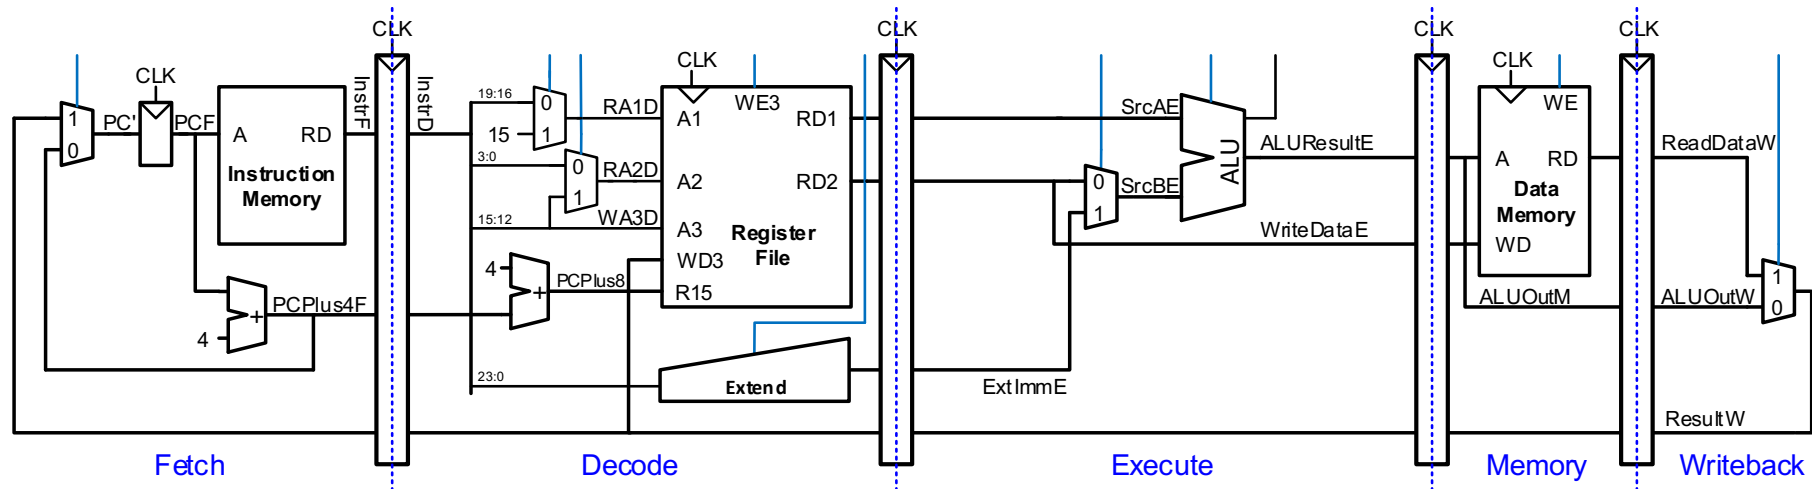
- ❑ Stages and their boundaries are indicated in blue
- ❑ Signals are given a suffix (F, D, E, M, or W) to indicate the stage in which they reside

# Pipeline Operation

- Consider the example instruction sequence

I1:	ADD	R0,	R5,	#10
I2:	ADD	R1,	R5,	#10
I3:	ADD	R2,	R5,	#10
I4:	STR	R0,	[R7,	#4]
I5:	STR	R1,	[R7,	#8]
I6:	STR	R2,	[R7,	#12]

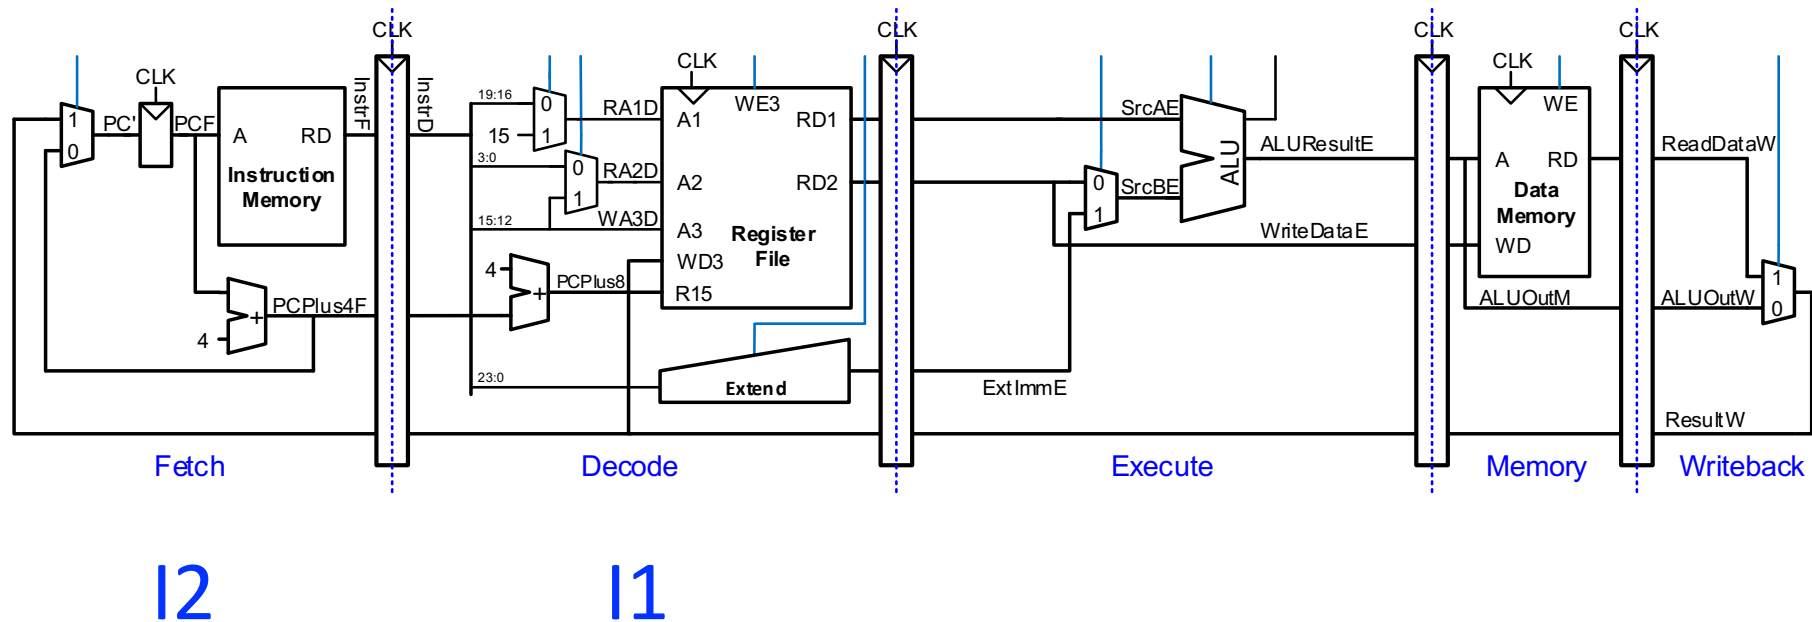
# Pipeline Operation: Cycle 1



1

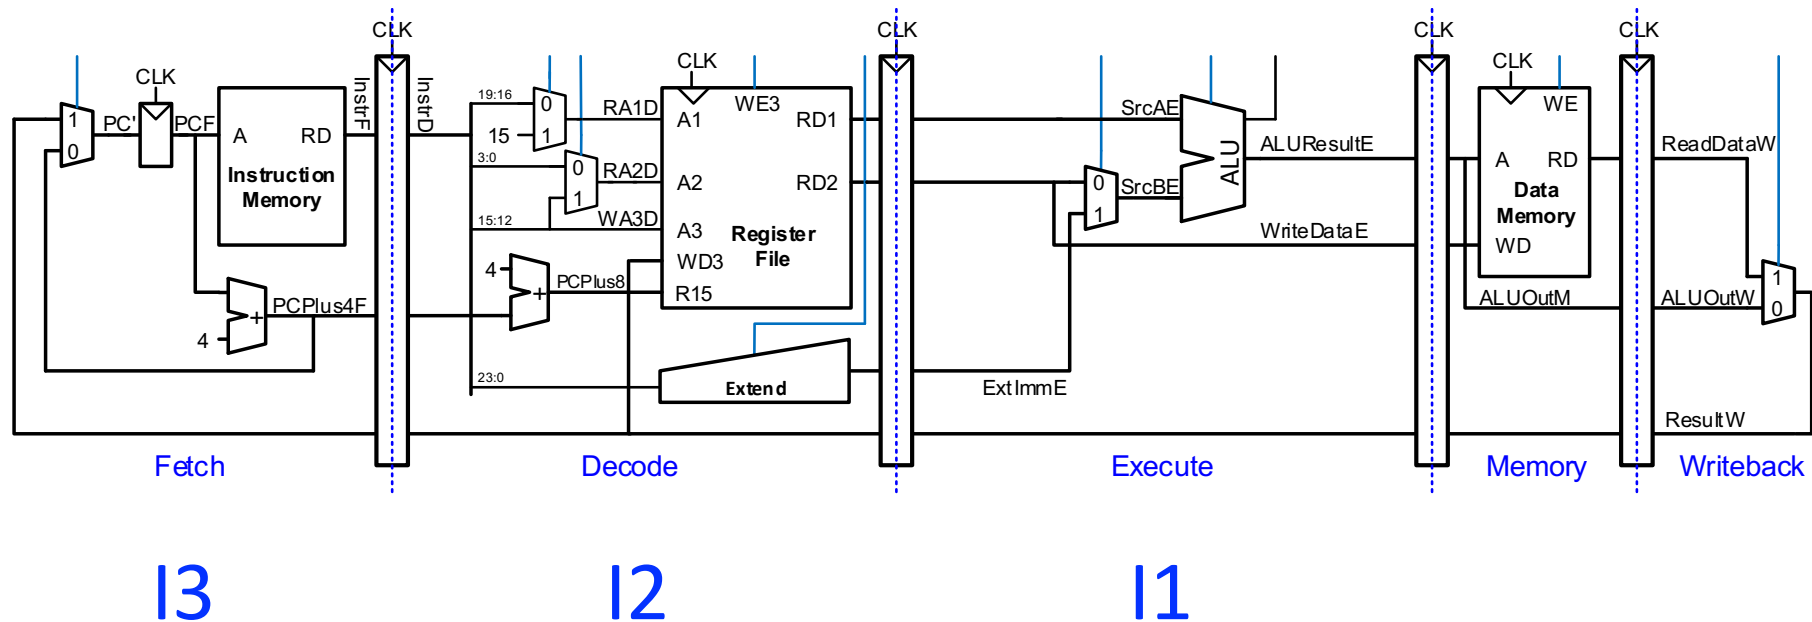
❑ Is the pipeline fully utilized? **NO**

# Pipeline Operation: Cycle 2



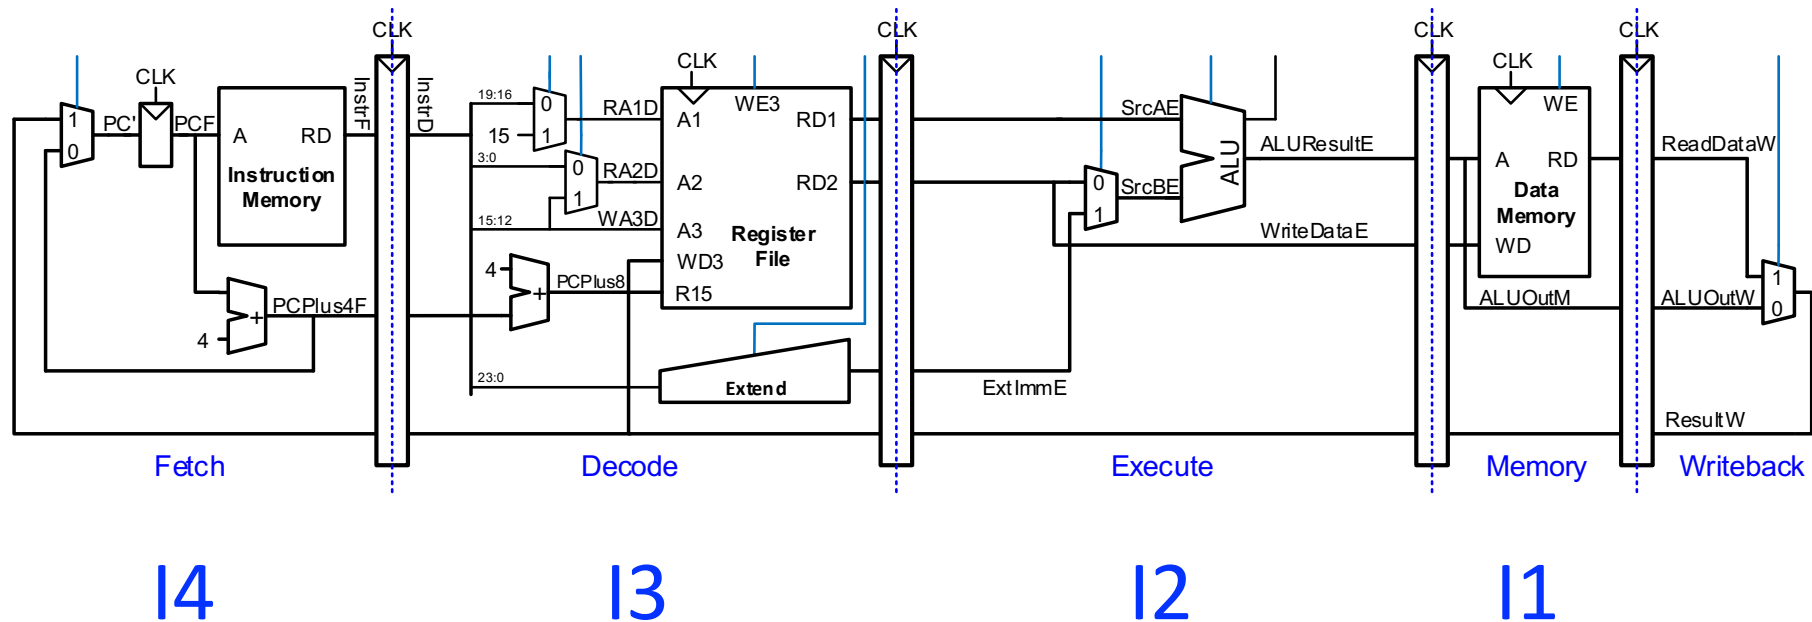
❑ Is the pipeline fully utilized? **NO**

# Pipeline Operation: Cycle 3



❑ Is the pipeline fully utilized? **NO**

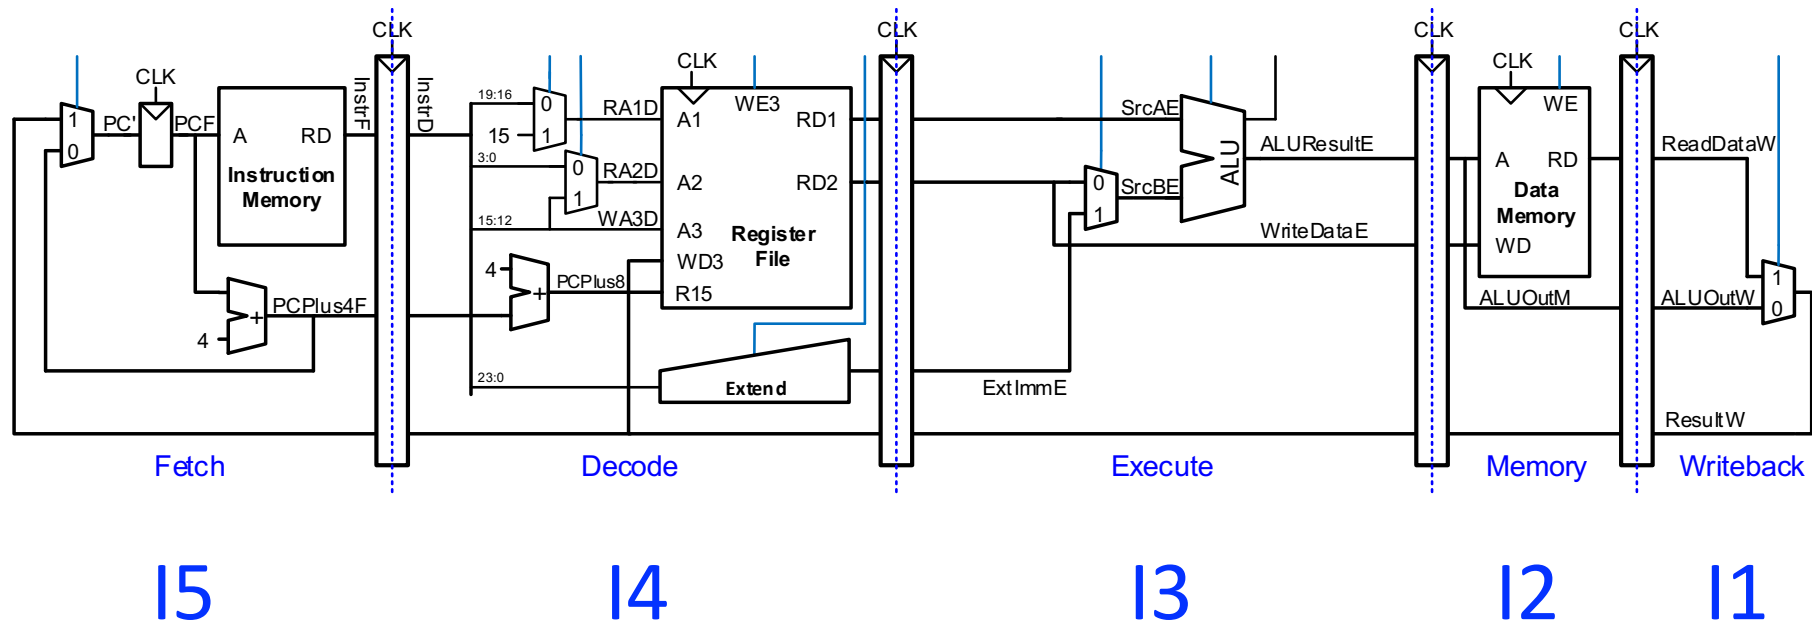
# Pipeline Operation: Cycle 4



❑ Is the pipeline fully utilized? **NO**

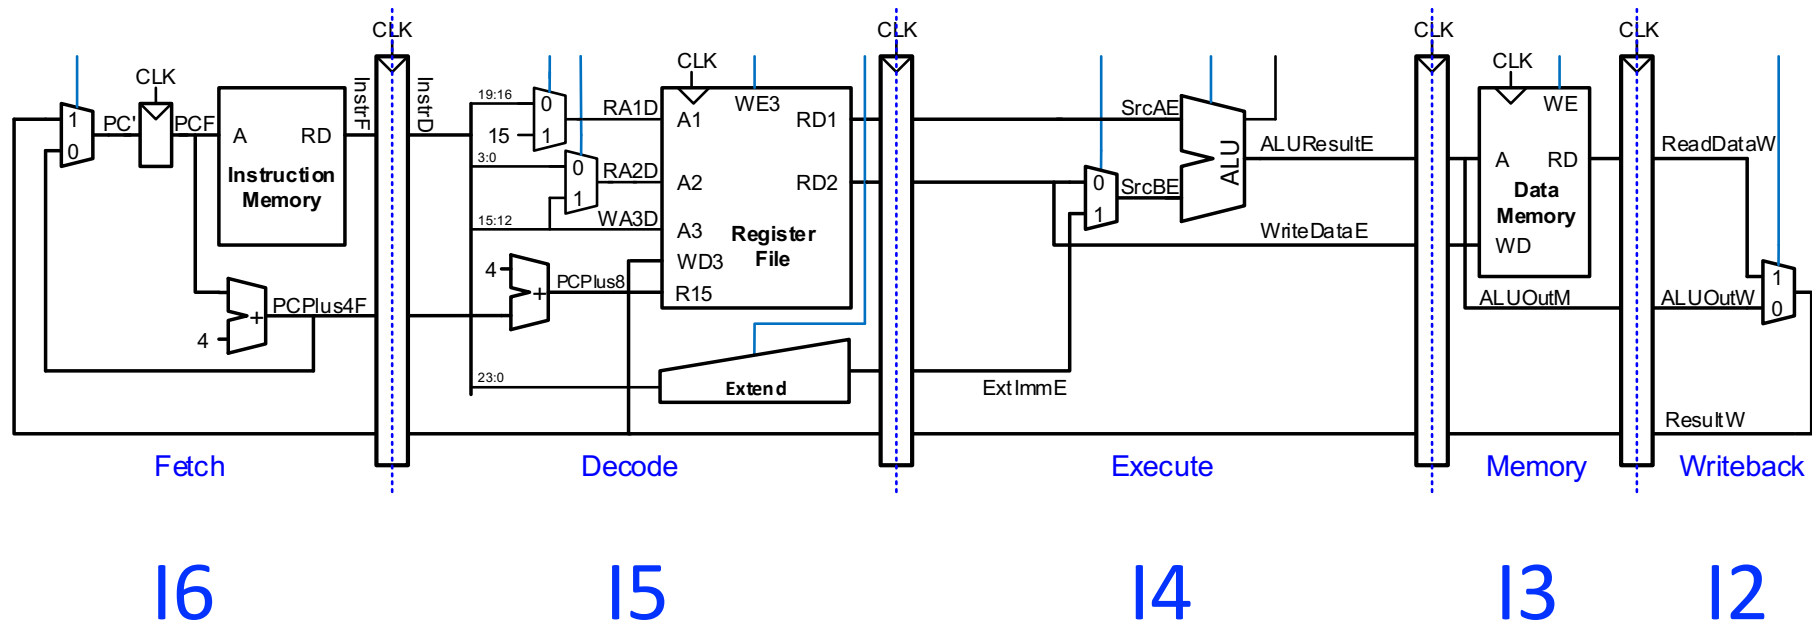


# Pipeline Operation: Cycle 5



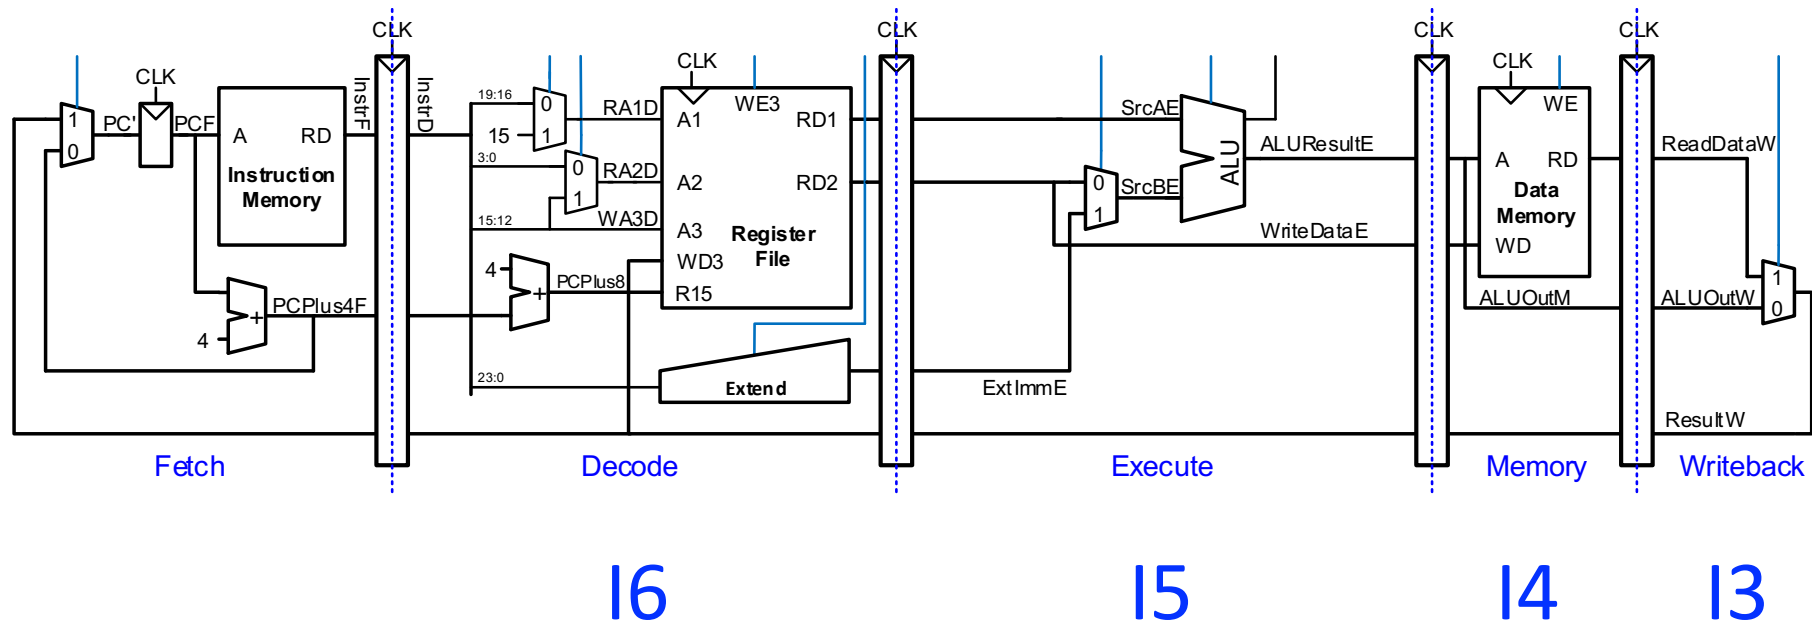
☐ Is the pipeline fully utilized? YES

# Pipeline Operation: Cycle 6



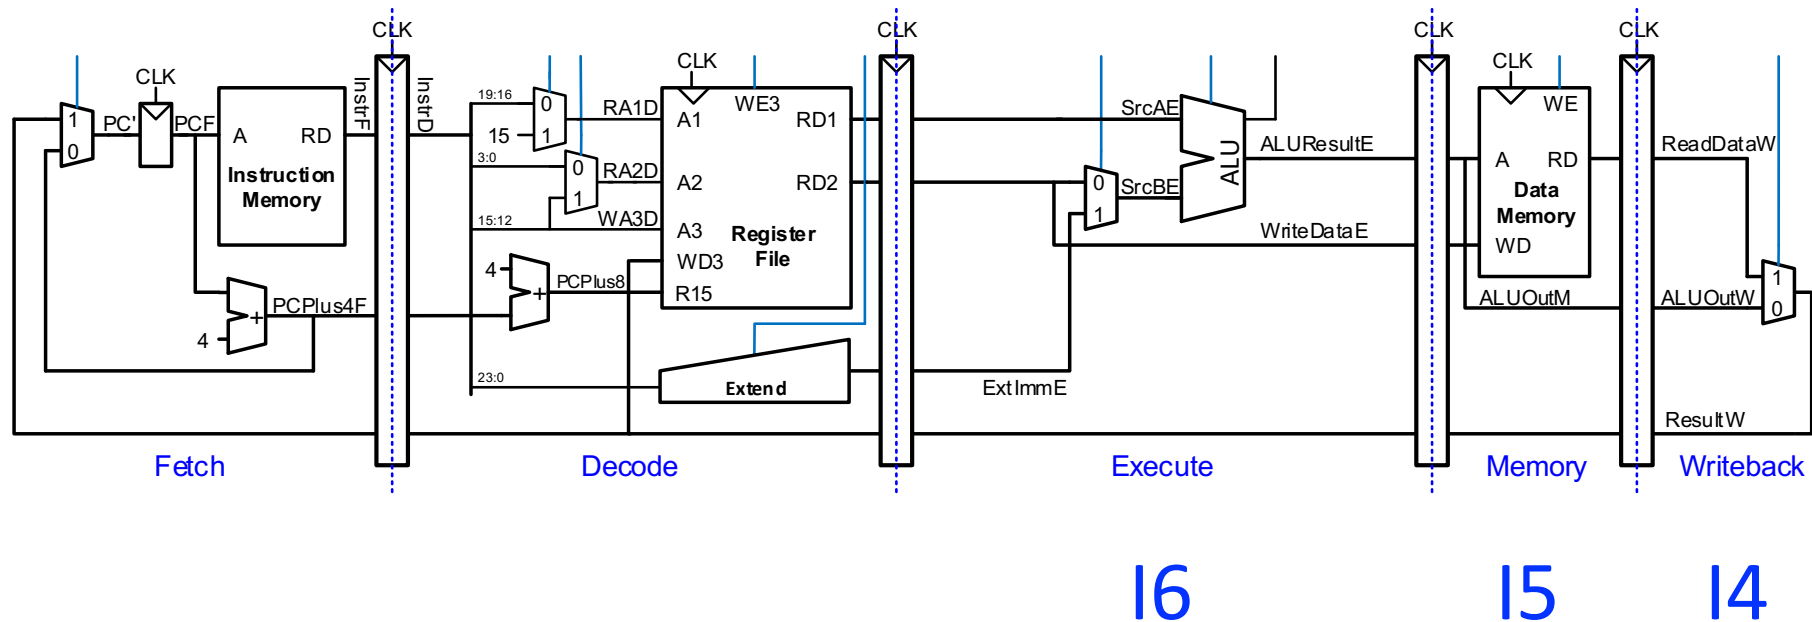
☐ Is the pipeline fully utilized? YES

# Pipeline Operation: Cycle 7



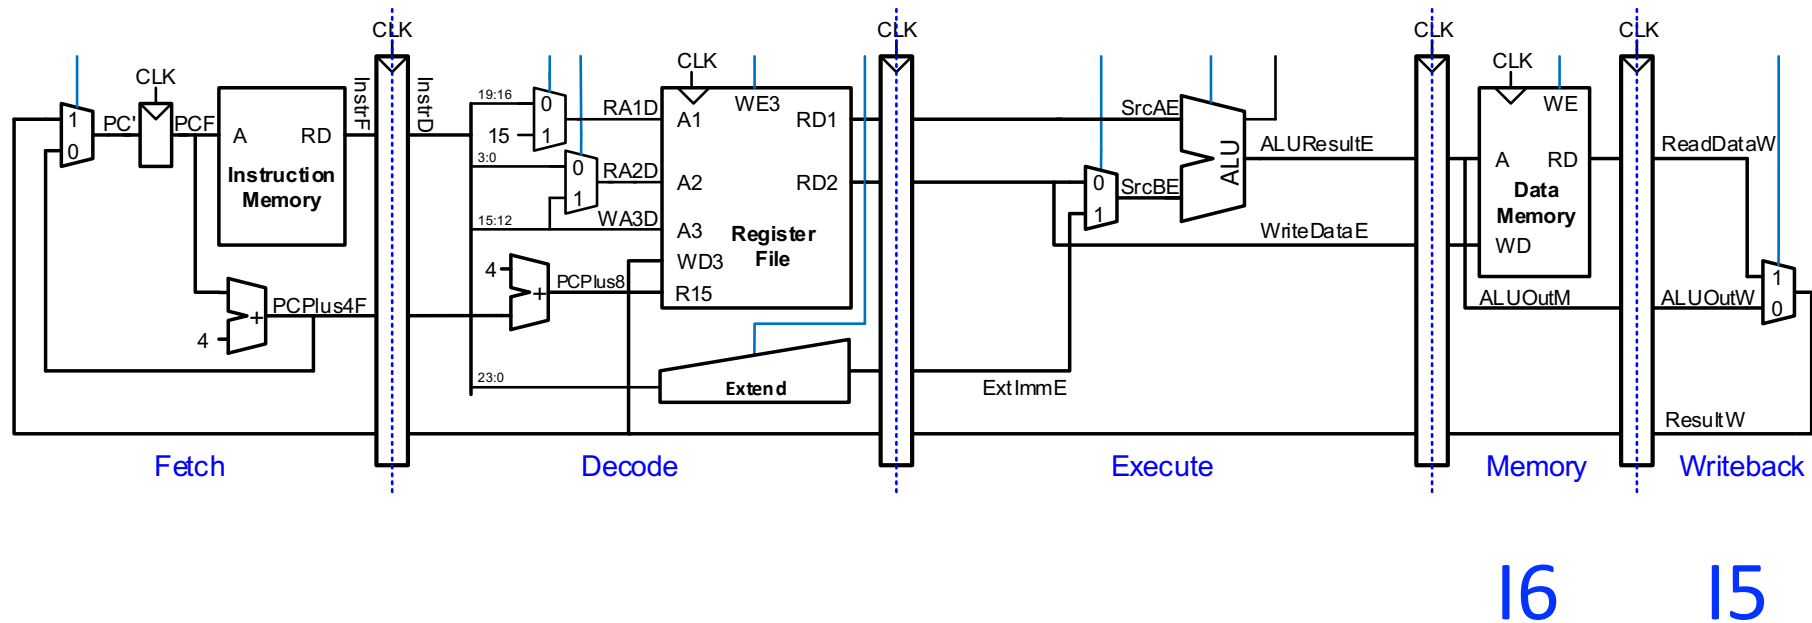
❑ Is the pipeline fully utilized? **NO**

# Pipeline Operation: Cycle 8



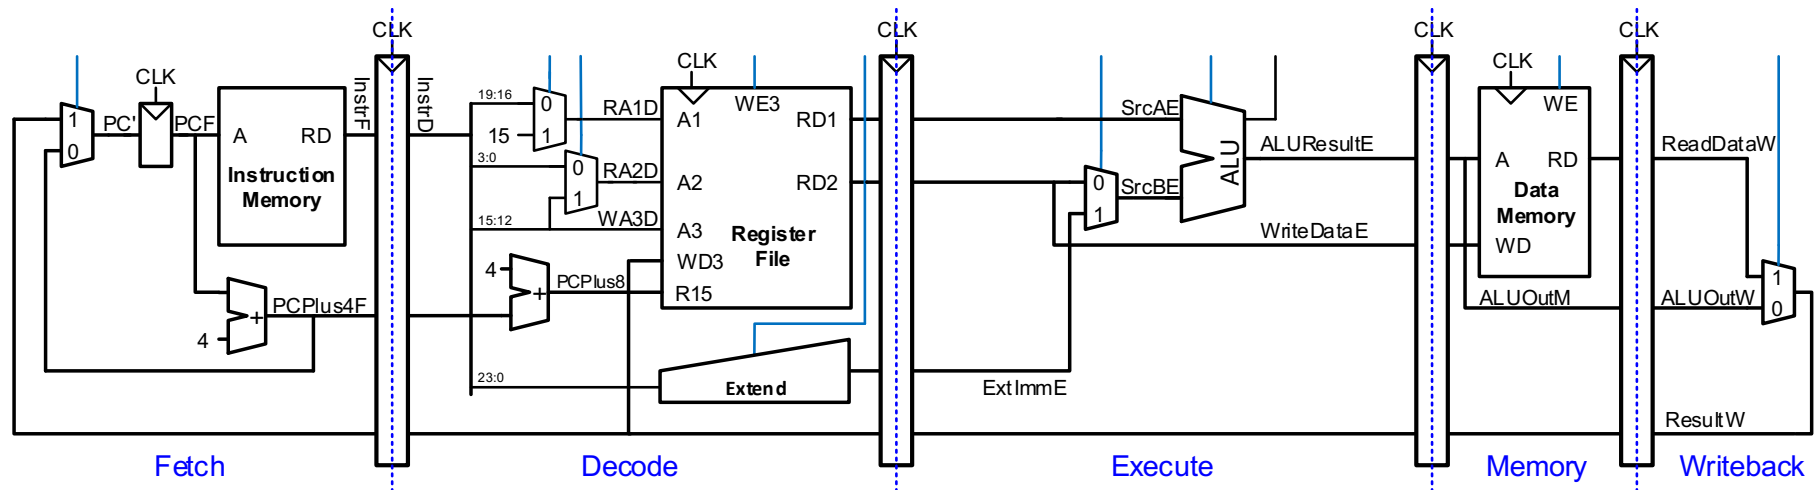
❑ Is the pipeline fully utilized? **NO**

# Pipeline Operation: Cycle 9



❑ Is the pipeline fully utilized? **NO**

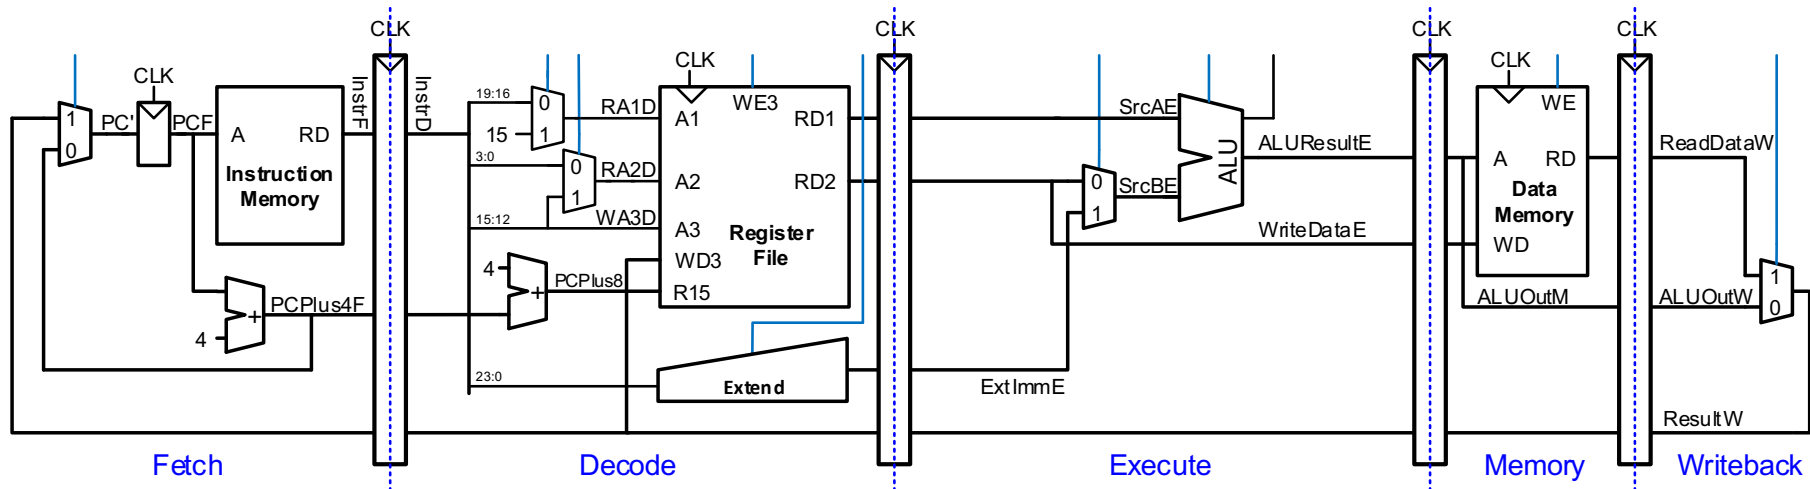
# Pipeline Operation: Cycle 10



16

❑ Is the pipeline fully utilized? **NO**

# Pipeline Operation




❑ No more instructions to execute

# Performance Analysis

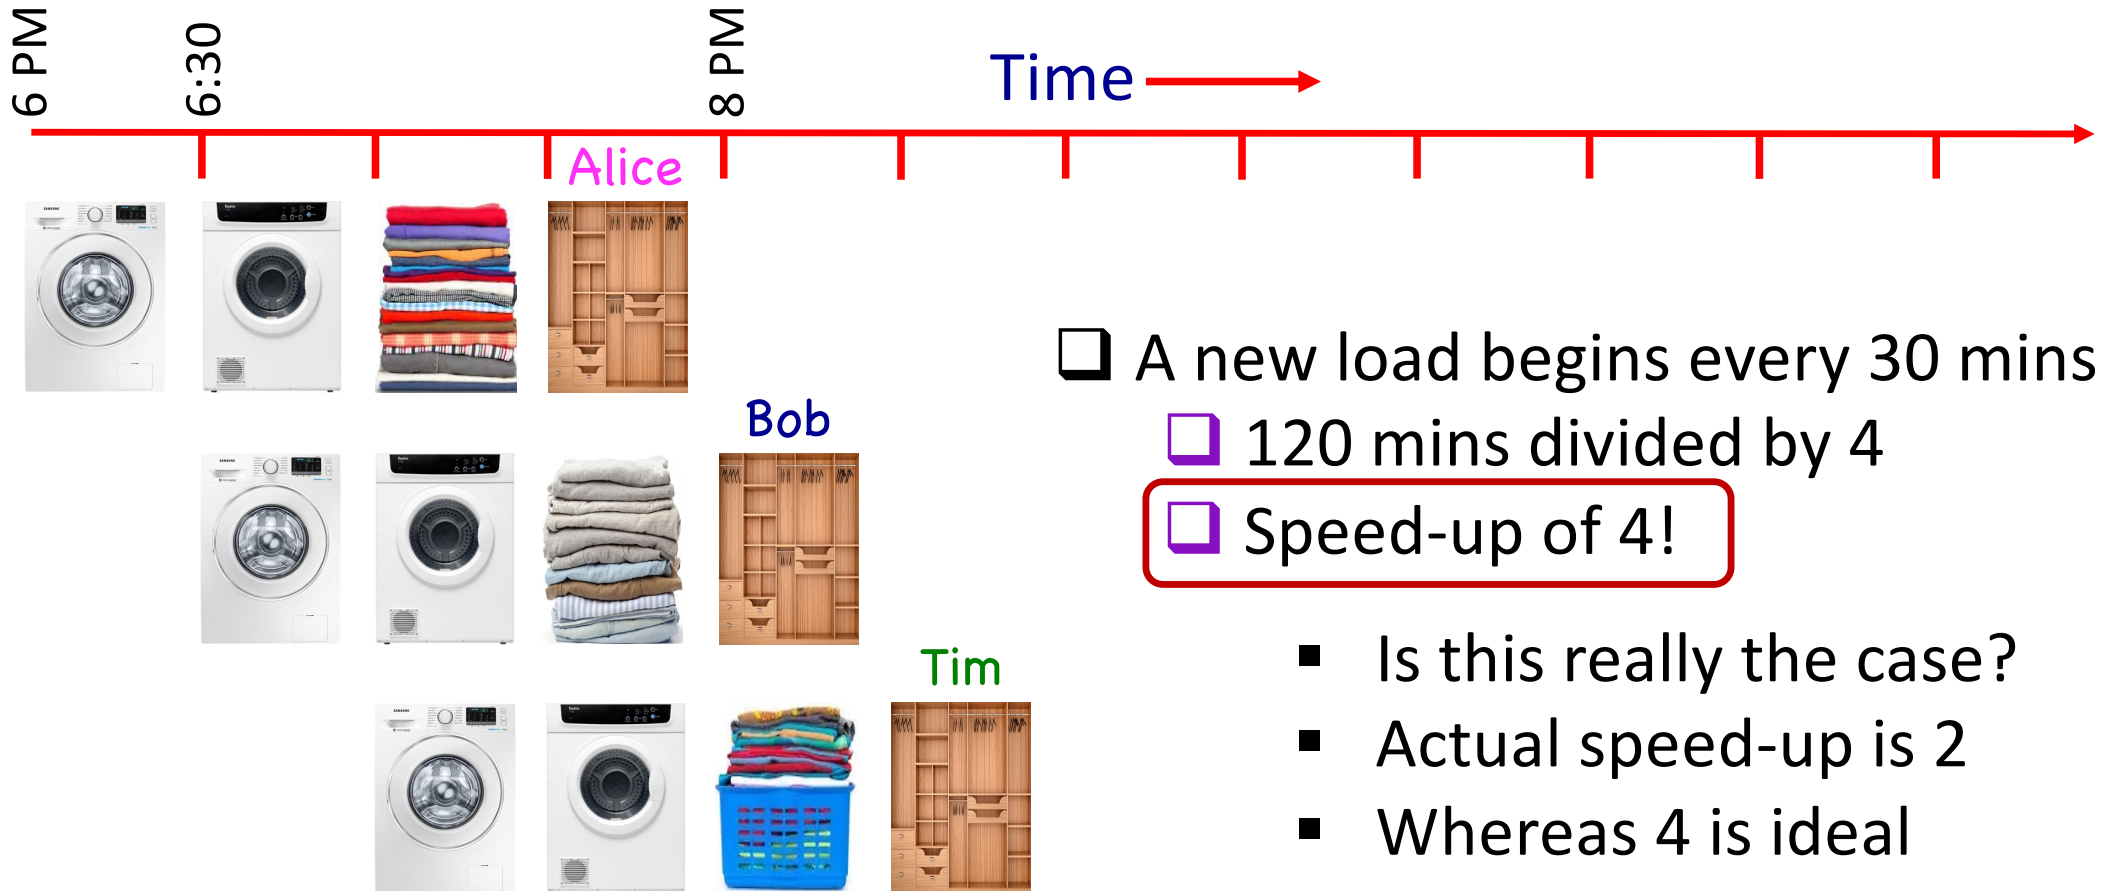
- The **6 instructions** took **10 cycles** to finish execution
- Cycles per Instruction (CPI) is :  $10/6 = 1.66$ 
  - Conversely, instruction per cycle (IPC) is: 0.6
- **Ideally**, we want the IPC to be close to 1
  - **One instruction finished every cycle**
- Why is IPC **less than 1**?
  - It takes some time to fill and some time to drain the pipeline
    - During this time **pipeline is operating below its potential**



# Sequential Laundry



# Recall: Pipelined Laundry



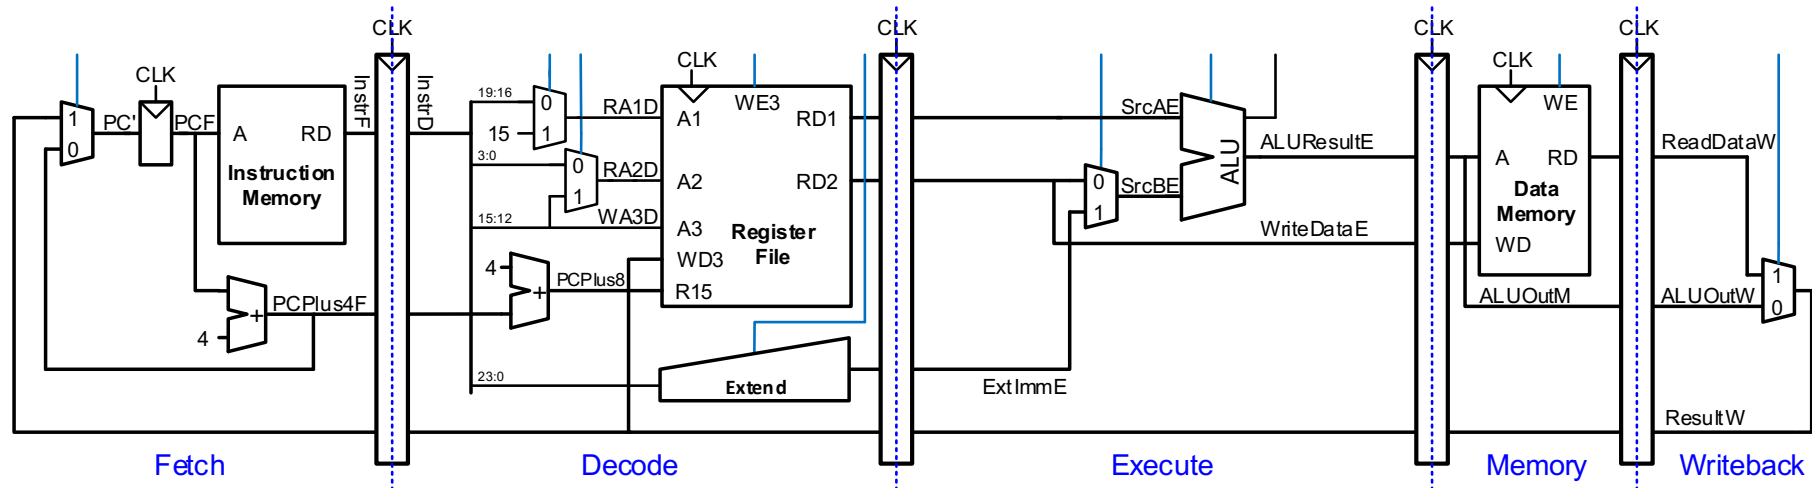
# Pipeline Idealism vs. Reality

- **Pipeline fill time:** The time it takes to fill the pipeline and make it operate at maximum efficiency
- **Pipeline drain time:** The time that is wasted when there is no more work to do in the pipeline
- The two factors limit the pipeline from delivering ideal speed-up
  - In the case when the amount of work is small relative to the number of stages in the pipeline
- Let's revisit the previous example

# Performance Analysis

- The **6 instructions** took **10 cycles** to finish execution
- Cycles per Instruction (CPI) is :  $10/6 = 1.66$ 
  - Conversely, instruction per cycle (IPC) is: 0.6
- What if we have 1 billion instructions instead of 6?
  - $CPI = (4 + 1000000000)/1000000000 = \sim 1$
- Computer programs **execute billions of instructions**, so the overhead of filling/draining is amortized

# Bug in Pipelined Hardware!



- There is a “hardware bug” in the pipelined microarchitecture
  - Can you spot it?